



# The Memory Hierarchy

Registers, caching, RAM, disks, swapping, networks...



# The Memory Bottleneck

Registers are very, very fast:

Can be read/written as part of a single CPU cycle.

Main memory is up to **100x slower!**

Consider this code for doubling the value stored at address 1A0:

Load 1A0

CPU has to wait with next instruction until memory has sent the data

Add 1A0

But we're using the data from 1A0 multiple times!

Store 1A0

# Solution: Caching

A **cache** sits between the registers and the RAM:

- Additional memory in the CPU
- Faster than RAM, slower than registers
- Smaller than RAM, larger than registers
- Keeps recently loaded values close to the CPU

# Caching: Implementation

- Fully implemented in hardware
- “Transparent” for the program:

`Load 1A0` loads from cache if possible, from RAM otherwise!

`Store 1A0` may store only to cache (for later transfer to RAM), or both to cache and RAM, or only to RAM

# Caching tricks

- Programs often access consecutive addresses, e.g.
  - searching through some text from start to finish
  - making all pixels in an image brighter
  - playing video or audio data
- When you `Load` a value that's not yet in the cache, the cache will already load the next few values as well
- When cache is full, the “oldest” values are thrown out
  - Less likely to need those again

# What makes caching tricky

- “Cache misses” can make a program go 100x slower!
- Performance of code
  - may be difficult to predict
  - may depend heavily on CPU cache architecture
- High-performance code may have to be written with cache in mind
- Remember: programs themselves are stored in RAM (and therefore also cached)

# Dealing with Big Data

What to do when you don't have enough RAM?

## Swapping

- Write unused parts of RAM to disk
- Read back when required
- Almost opposite of caching
- Implemented by combination of hardware and operating system (*virtual memory*, more later)

# Modern Memory Hierarchy

| Memory         | Speed                          | Size           | Connection to CPU      |
|----------------|--------------------------------|----------------|------------------------|
| Registers      | Very fast                      | A few words    | Inside CPU             |
| L1 Cache       | Fast (3-5x slower)             | ~100 KByte     | Inside CPU             |
| L2, L3 Cache   | Fast (25-50x slower)           | ~1-8MBytes     | Inside or close to CPU |
| RAM            | Medium (50-100x slower)        | Several GBytes | Main bus               |
| Hard Disk, SSD | Slow (1000x slower)            | Terabytes      | External (SATA, USB)   |
| Network        | Very slow (10000x-infinity...) | ?              | External (PCIe, USB)   |



# Summary

- Fast memory access is critical for code performance
- Accessing RAM is up to 100x slower than registers
- Cache keeps values around for fast access
  - Recently loaded values
  - Values close in RAM to recently used values
- Modern processors unthinkable without cache
- Swapping is opposite of caching: use disk to simulate more RAM than available

EOF